

DSI Claude API Credit Award — Laura Review Memo

What is ready

The scientific narrative and workload model are prepared. Submission still requires Laura’s review and CV, confirmation of the external-grant details and team roles, completion of any required PI-specific fields, and final portal entry and approval.

1. Portal-ready copy/paste application

`docs/grants/2026-09-dsi-claude-portal-ready.md`

Contains the title, abstract, project description, purpose of Claude credits, CAD \$75,000 request, detailed budget justification, AI-safety response, team text, and external-grant placeholder. The portal fields are deliberately plain text with no equations, tables, or citation-dependent formatting.

2. Application packet / checklist

`docs/grants/2026-09-dsi-claude-application-packet.md`

Includes the portal content plus administrative notes, PI items, and final submission checklist.

3. Complete technical proposal

`docs/grants/2026-09-dsi-anthropic-ai-privacy-firewall-proposal.pdf`

Full methodological proposal for review, including the statistical framework, Jev/Laya architecture, benchmark and data-source strategy, uncertainty/calibration methods, harm-weighted evaluation, phased stop/go gates, training plan, software outputs, publication path, and potential public-health pilot.

4. Editable technical source

`docs/grants/2026-09-dsi-anthropic-ai-privacy-firewall-proposal.md`

Current project framing

Title: Measure Before You Share: Calibration, Uncertainty, and Dual-Gate Safety for AI Access to Medical and Public-Health Data

Request: CAD \$75,000 in Claude API credits over 12 months.

Core scientific question: Can calibrated statistical and semantic evidence determine whether an original, transformed, synthetic, or summarized representation of medical/public-health data is appropriate to expose to a cloud AI agent for a specific task?

Main contributions: - privacy-aware AI data-access decisions rather than generic PII detection; - interpretable statistical + semantic typed-decision models; - calibration, uncertainty, abstention, correlated-error and harm-weighted evaluation; - reusable training/adaptation and classifier-evaluation infrastructure in Which; - privacy implementation in DataGangerR; - 3–5 supervised Biostatistics/data-science trainees; - open benchmark, open-source software, methods manuscript, and optional partner pilot; - phased stop/go design so failure of a later modelling stage does not cause project failure.

Questions for Laura

1. Budget level

We currently request **CAD \$75,000**. The bottom-up plan projects approximately CAD \$73,660 of usage, leaving about CAD \$1,340 for workload and exchange-rate variation. The estimate is based on explicit token totals and a workload decomposition in `docs/grants/2026-09-dsi-claude-budget-model.md`.

Question: Does CAD \$75,000 feel appropriate for this competition, or would you prefer a lower request? If you know whether DSI commonly offers a smaller allocation than requested, that would help us decide whether to keep the current amount or explicitly state a reduced-award contingency.

If a smaller amount is awarded, our proposed priority order is: 1. benchmark + core statistical evaluation; 2. open-source DataGangerR/Which framework; 3. trainee research; 4. robustness/adversarial replication; 5. model-version re-evaluation; 6. optional pilot.

2. External grant evidence

The current placeholder is:

- CIHR / Diabetes Canada — Embracing Diversity to Achieve Precision & Health Equity Team Grant
- *Developing a Precision Participatory and Multi-Level Approach for Population-Based Diabetes Risk Assessment to Address Inequities in Type 2 Diabetes*
- Funded, 2025
- Laura as lead/PI

Question: Please confirm the exact agency/program, title, status, and lead-applicant wording we should enter in the portal.

3. Team/portal role for Lennon

The application currently describes Lennon as **project lead / collaborator / software-methods lead**, not PI.

Question: When you open the Good Grants form, is there an appropriate collaborator/co-investigator role for Lennon, or does the portal only accept PIs and trainees?

4. Trainee mechanism

We propose **3–5 Biostatistics/data-science trainees**, potentially through practicum or supervised research-project mechanisms.

Question: Is this a realistic mechanism from your perspective? Should we keep the trainees unnamed at submission or identify likely students/supervision routes now?

5. Data/privacy partners and external validation

The primary benchmark is feasible using public sources and synthetic/counterfactual cases, so the study does not depend on institutional data access. We also propose seeking abstracted/de-identified decision cases or expert adjudication from privacy/data-access organizations such as DSI, Tri-Agency programs, ICES, DLSPH units, PHO, or comparable partners.

Questions: - Are you comfortable naming these as potential partners without implying commitment? - Is there a privacy/data-governance group or individual you think would be particularly appropriate for benchmark adjudication or external validation? - Would you prefer us to approach any partner before submission, or keep this as a post-award activity?

6. Public-health pilot

The proposal currently says “**potentially with PHO, subject to organizational approval.**”

Question: Should we keep PHO named as a potential pilot site, or use the more generic “public-health partner” wording until there is formal interest?

7. Gold-standard adjudication

The benchmark requires a smaller rigorously human-adjudicated core and an independently double-reviewed high-value subset.

Question: Who would you view as the right reviewer mix—biostatistics/data science plus privacy/data governance, or should we add another type of expertise?

8. Full-proposal attachment

The published call appears to require portal text plus PI CVs, with no separate full-proposal upload listed.

Question: When you access the live portal, do you see any optional proposal/attachment field beyond the CV upload? If yes, we can provide the full proposal PDF directly.

9. PI materials and submission

Please confirm: - PDF CV to use; - demographic survey status, if required; - any institutional/unit wording that should be changed; - whether you would like us to prepare the final portal entries exactly as they should be pasted before you submit.

What we need from Laura to finalize

The minimum items needed are: 1. review/approval of the scientific framing; 2. decision on the CAD \$75,000 request; 3. exact external-grant information; 4. CV PDF; 5. portal role guidance for Lennon; 6. preference on naming PHO/other potential partners; 7. confirmation of trainee framing; 8. confirmation whether the portal permits a full-proposal attachment.

Once these are resolved, the remaining work is primarily copy/paste, final word-count verification, and submission.